

TENSOR PROGRAMS V

# Zero-Shot Hyperparameter Transfer

Tuning Large Neural Networks via  
Maximal Update Parametrization

 Greg Yang, Edward J. Hu, et al.

 Microsoft & OpenAI

# Contents

01

## The HP Tuning Challenge

Why hyperparameter tuning becomes prohibitively expensive as models scale to billions of parameters

02

## Standard Parametrization Problems

Mathematical analysis of why HPs don't transfer in conventional neural network implementations

03

## Maximal Update Parametrization

The theoretical foundation of  $\mu$ P and how it enables zero-shot HP transfer across model scales

04

## $\mu$ Transfer Algorithm

The complete procedure: from small proxy tuning to large model deployment with massive cost savings

05

## Experimental Validation

Large-scale results on BERT and GPT-3 demonstrating superior performance with dramatically reduced tuning costs

06

## Practical Implementation

Key properties, theoretical insights, and guidelines for implementing  $\mu$ Transfer in practice

# The Hyperparameter Tuning Challenge

## Why HP Tuning is Critical

### Poor HPs → Subpar Performance

Bad hyperparameters cause training instability and reduced model quality

### Baseline Comparability Issues

Varying HP tuning makes published results hard to compare fairly

### Scaling Exacerbates the Problem

Billions of parameters make tuning prohibitively expensive

## The Problem: Optimal Learning Rate Shifts with Width (Standard Parametrization)




Figure 1 (Left): In standard parametrization, optimal learning rate shifts by orders of magnitude as Transformer width increases. Using small model HPs on large models causes divergence.

## The $\mu$ Transfer Solution

**Core Idea:** Tune HPs on a small proxy model, then zero-shot transfer them to the full-size target model without direct tuning.

13M → 350M

**BERT-large**

Outperformed published numbers

40M → 6.7B

**GPT-3**

Only 7% tuning cost

Key Innovation

$\mu$ P

Enables HP stability

# Why Parametrization Matters: A Mathematical Primer

## Central Limit Theorem Analogy

The Central Limit Theorem (CLT) provides the foundation for understanding correct parametrization in neural networks.

Given iid samples  $x_1, \dots, x_n \sim N(0, 1)$ :

$$\frac{1}{\sqrt{n}} (x_1 + \dots + x_n) \rightarrow N(0, 1)$$

The scaling factor  $\frac{1}{\sqrt{n}}$  is the "correct" order for convergence.

$$c_n = \frac{1}{n}$$

$$\rightarrow 0$$

Shrinks to zero

$$c_n = 1$$

$$\rightarrow \infty$$

Blows up

$$\begin{aligned} c_n &= \frac{1}{\sqrt{n}} \\ &\rightarrow N(0, 1) \\ &\text{Converges} \end{aligned}$$

## Extension to Multiple HPs

For multiple hyperparameters  $c_1, \dots, c_k$ :

$$F_n(c_1, \dots, c_k) = E[f((c_1 + \dots + c_k)(x_1 + \dots + x_n))]$$

The correct reparametrization is:

$$\alpha_i = c_i \sqrt{n}$$

**Critical Insight:** All HPs must be correctly parametrized. Fixing some while neglecting others causes compensation distortions.

## Connection to Neural Networks

Neural network training parallels this optimization problem. Correct parametrization enables direct HP transfer from narrow to wide networks.

### Standard Parametrization (SP)

- Initialization:  $\Theta(1/\text{fan\_in})$
- Learning rate:  $\Theta(1)$
- ✗ No well-defined infinite-width limit

### Maximal Update Parametrization ( $\mu P$ )

- Input:  $\Theta(1/\text{fan\_out})$
- Hidden:  $\Theta(1/\text{fan\_in})$
- Output:  $\Theta(1/\text{fan\_in}^2)$
- ✓ Enables feature learning

## Key Theoretical Result:

$\mu P$  is the unique parametrization that allows HP transfer across width while preserving feature learning in the infinite-width limit.

## Core Principle

If we parametrize correctly, optimal HPs from a narrow network transfer approximately to a wide network.

# Hyperparameter Instability in Standard Parametrization

MLP on CIFAR-10: Optimal Learning Rate Shifts





Figure 3: MLP with 2 hidden layers, ReLU activation, SGD optimizer. Optimal LR shifts  $-4x$  (2 orders of magnitude in log space) as width increases from 256 to 8192.

Optimal LR  
-12 → -6

Transformer on Wikitext-2: Same Instability



## The Consequences

1

### Divergence Risk

Using small model optimal LR on large model causes divergence or poor performance.

2

### No Transfer Guarantee

HP instability means each model size requires expensive independent tuning.

3

### Beyond Learning Rate

Initialization scale and other HPs also shift with width (Fig. 18).

## Experimental Setup Details

### MLP

- 2 hidden layers
- Width: 256 → 8192
- Activation: ReLU
- Loss: Cross-entropy

### Transformer

- d\_model: 128 → 8192
- Depth: 2
- Optimizer: Adam
- Dataset: Wikitext-2

# Maximal Update Parametrization ( $\mu$ P)

## What is $\mu$ P?

The Maximal Update Parametrization ( $\mu$ P) is a principled scaling of initialization and learning rates that ensures each layer is updated on the same order during training, regardless of width.

### Key Design Principle:

All hidden activations update at the same speed in terms of width. No layer updates too fast or too slow as network scales.

#### Standard Parametrization

- Layers update at different speeds
- Output blows up with width
- No HP transfer possible

#### Maximal Update Parametrization

- Uniform update speed across layers
- Stable activations with width
- Enables zero-shot HP transfer

## $\mu$ P Scaling Rules (Table 3)

### Input Weights & All Biases

Init Var:  $1/\text{fan\_out}$  | Adam LR: 1

### Output Weights

Init Var:  $1/\text{fan\_in}$  | Adam LR:  $1/\text{fan\_in}$

### Hidden Weights

Init Var:  $1/\text{fan\_in}$  | Adam LR:  $1/\text{fan\_in}$

Transformer Addition: Attention logits scaled by  $1/d$  instead of  $1/\sqrt{d}$

## $\mu$ P Solves the Problem: Stable Optimal LR




Figure 1 (Right): In  $\mu$ P, optimal learning rate is stable across all widths. Best LR for width-128 is also best for width-8192. Performance improves monotonically with width.

## Why "Maximal Update"?

$\mu$ P enables "maximal" feature learning in the infinite-width limit:

- ✓ Each layer contributes meaningfully to updates
- ✓ Activations don't vanish or explode
- ✓ Preserves feature learning (unlike NTK limit)

**Contrast with NTK:** Neural Tangent Kernel parametrization loses feature learning ability in infinite width.  $\mu$ P preserves it, making it suitable for modern large-scale pretraining.

# How $\mu$ P Works: Defects of SP and $\mu$ P Fixes

## The Blow-Up Problem in Standard Parametrization

In SP, network output blows up with width after just **1 step of SGD** due to imbalanced layer updates.

### 1-Hidden-Layer Linear Perceptron Example:

$$f(x) = V^T U x$$

where  $U, V \in \mathbb{R}^{n \times 1}$  (width n)

SP Initialization:

$$V \sim N(0, 1/n), U \sim N(0, 1)$$

After 1 SGD step (LR=1):

$$f'(x) = (V^T U + \theta U^T U + \theta V^T V + \theta^2 U^T V)x$$

**The Problem:**

$$U^T U = \Theta(n) \text{ by Law of Large Numbers}$$

Output blows up linearly with width n

## $\mu$ P Prevents Blow-Up

Same network in  $\mu$ P:

$\mu$ P Initialization & LR:

$$V \sim N(0, 1/n^2), U \sim N(0, 1)$$

$$\eta_V = 1/n, \eta_U = n$$

## Empirical Evidence: Transformer Activations




Figure 5: Standard deviation of coordinate changes from initialization over 4 Adam steps. Logits and attention logits blow up in SP but remain stable in  $\mu$ P.

## Why Scaling Down LR in SP Fails

Simply reducing learning rate with width doesn't work:

- ✗ **Word embeddings** update by width-independent amount per step. Small LR means they don't learn in large models.
- ✗ **Input vs output layers** have different scaling. SP updates them at different rates.

# Which Hyperparameters Can Be μTransferred?

## Three Categories of Hyperparameters

### Category 1: Transferable

These HPs can be transferred from small to large models:

- Learning rate
- Momentum
- Adam  $\beta_1, \beta_2$
- LR schedules
- Initialization variance
- Parameter multipliers
- Per-layer HPs

### Category 2: Not Transferable

Regularization HPs — their purpose depends on both model and data size:


- Dropout probability
- Weight decay
- Data augmentation
- Label smoothing

### Category 3: Transferred Across

Scale-defining HPs across which others are transferred:

- Width
- Depth
- Batch size
- Sequence length
- Training steps

## Empirical Validation: HP Stability Across Width & Depth



**Figure 4:** Four representative HPs on pre-LN Transformers in μP: learning rate, output weight multiplier  $\alpha_{\text{output}}$ , initialization std  $\sigma$ , and LR schedules (constant, linear, cosine, etc.). Models trained on Wikitext-2 for 10k steps.

## Transfer Requirements

For stable transfer, ensure:

- Width  $\geq 256$
- Depth  $\geq 4$
- Batch size  $\geq 32$
- Seq length  $\geq 128$

## Limitations

- Init std doesn't transfer well across depth
- Depth transfer only works for pre-LN
- Optimum can shift slightly at scale
- But shift has small impact vs SP

# The $\mu$ Transfer Algorithm

## Algorithm 1: $\mu$ Transfer

### 1 Parametrize Target Model

Parametrize the target (large) model in Maximal Update Parameterization ( $\mu$ P)

### 2 Tune Proxy Model

Tune a smaller version (in width and/or depth) of the target model

### 3 Copy Hyperparameters

Copy tuned hyperparameters to the target model

#### Result:

Near-optimal HPs on the full-sized target model **without directly tuning it at all!**

## Key Benefits

- ⚡ **Massive Speedup:** Tuning cost independent of target model size

- ✍ **Tune Once:** Single proxy tune serves whole model family

- ↪ **Better Performance:** Outperforms SP with optimal LR

## Large-Scale Results: BERT & GPT-3

### BERT (350M params)

Proxy: 13M model

Tuning cost = 1 BERT-large pretrain

#### Test Loss:

Megatron: 1.683

$\mu$ Transfer: 1.731 → Better

### GPT-3 (6.7B params)

Proxy: 40M model

Tuning cost: only 7% of pretraining

#### Validation Loss:


Original [7]: 2.03

$\mu$ Transfer: 1.98 → Better

## Revolutionary Efficiency:

For models like GPT-3, HP tuning is not feasible at all without  $\mu$ Transfer. The approach reduces tuning cost from 100% to 7% of pretraining.

## Compute-Performance Pareto Frontier



# Experimental Results: BERT and GPT-3

## BERT Pretraining Results

### Setup:

- Proxy: 13M params
- Target: BERT-base (110M) & BERT-large (350M)
- Tuning: 256 HP samples
- Training:  $10^5$  steps
- Cost  $\approx 1$  BERT-large pretrain
- No direct tuning of targets

### BERT-base Results:

Method	Test Loss	MNLI (m/mm)
<b>Megatron</b>	1.995	84.2/84.2
<b>Naive</b>	diverged	—
<b><math>\mu</math>Transfer</b>	<b>1.970</b>	<b>84.3/84.8</b>

### BERT-large Results:

Method	Test Loss	MNLI (m/mm)
<b>Megatron</b>	1.683	86.3/86.2
<b>Naive</b>	diverged	—
<b><math>\mu</math>Transfer</b>	<b>1.731</b>	<b>87.0/86.5</b>

## GPT-3 6.7B Results

### Setup:

- Target: 6.7B model
- Proxy: 40M model
- Size ratio: 168×
- Tuning cost: 7% of pretrain
- Used FP32 (vs FP16 baseline)
- Relative attention

### Key Evaluation Results:

Validation Loss	<b>1.98 &lt; 2.03</b>
PTB Perplexity	<b>11.4 &lt; 13.0</b>
LAMBADA Zero-Shot	<b>73.5% &gt; 70.8%</b>
HellaSwag Zero-Shot	<b>72.0% &gt; 66.7%</b>

### Remarkable Achievement:

$\mu$ Transferred 6.7B model performs comparably to the twice-as-large 13B model from [7], despite using FP32 and encountering numerical issues.

## Practical Impact

- ⌚ **Cost Reduction:** From prohibitive to 7% of pretraining cost
- ⌚ **Time Savings:** 220× speedup for BERT-large tuning
- 🏆 **Performance:** Matched or exceeded published baselines

# Key Properties and Theoretical Insights

## Property 1: "Wider is Better" Throughout Training

In  $\mu$ P, wider models consistently achieve better training loss at any point, assuming output layer is zero-initialize d.

This contrasts sharply with SP:

- Standard Parametrization
- Curves cross frequently
- Wider  $\neq$  better reliably
- Scaling unpredictable

### Maximal Update Parametrization

- No curve crossing
- Wider strictly better
- Reliable scaling

## Stress Test: GPT-3 Transformer

Scaled width from 256 to 32,768 with fixed HPs:

- ✓ Wider models consistently matched or outperformed narrower ones
- ↳ Suggests wider models are strictly more data-efficient when scaled appropriately
- 🛠️ Provides cheap way to debug  $\mu$ P implementation: check "wider-is-better" early in training

## Practical Implications

### 1 Reliable Model Scaling

Researchers can confidently scale up models without fear of performance degradation from parametrization issues.

### 2 Better Compute Utilization

Small model tuning on individual GPUs increases parallelism and better utilizes organizational compute clusters.

### 3 Painless Exploration → Scaling

## Property 2: A Theoretical Puzzle

The existence of useful HP transfer reveals a deep theoretical question about neural net work training dynamics.

### The Paradox:

For useful transfer, we need the HP optimum to converge quickly with width, but the network fu nction should converge slowly.

#### HP Optimum

Converges at small width  
"Macroscopic" variable

#### Network Function

Converges very slowly  
"Microscopic" detail

### The Open Question:

Why does this separation exist? Where else should we expect useful HP transfer? T heoretically, it's unclear why HP optimum should converge faster than the function. This remains an open question for future theoretical work.

## Future Research Directions

➲ **Depth Parametrization:** Improve depth transfer, especially for post-LN Transformers

➲ **Finite-Width Corrections:** Fix slight optimum shifts at smaller scales

➲ **Regularization Transfer:** Transfer regularization HPs as function of model and data si ze

# Impact and Future Directions

## Transformative Impact

$\mu$ Transfer fundamentally changes how we approach hyperparameter tuning for large neural networks, making the previously impossible task of tuning billion-parameter models feasible and efficient.

### 🔑 Massive Speedup

Tuning cost becomes independent of target model size. From 100% to 7% of pretraining cost for GPT-3 6.7B.

### ↗ Better Performance

$\mu$ Transferred models outperform SP counterparts even with optimal tuning, due to balanced layer updates.

### 💡 Central Message

$\mu$ Transfer provides the first major practical payoff of infinite-width theory, transforming HP tuning from an art to a systematic, principled process.

### 💡 Tune Once for All

Single small proxy model provides HPs for entire model families (BERT, GPT, etc.).

### 💡 Better Utilization

Small model tuning on individual GPUs increases parallelism and cluster utilization.

## Future Research Directions

1

### Improving Depth Transfer

Fix initialization transfer across depth; extend to post-LN Transformers

2

### Finite-Width Corrections

Address slight optimum shifts at smaller scales via corrections to  $\mu$ P

3

### Regularization HP Transfer

Transfer regularization HPs as function of model and data size

4

### Theoretical Understanding

Explain theoretically why HP optimum converges faster than network function

## Practical Implementation

⟨⟩ PyTorch Package: `mup`  
`pip install mup`

☰ Documentation:  
[github.com/microsoft/mup](https://github.com/microsoft/mup)

The  $\mu$ Transfer algorithm and  $\mu$ P parametrization are now accessible to pr