

RLEF: Grounding Code LLMs in Execution Feedback

Reinforcement Learning for Iterative Code Synthesis with Multi-Turn Execution Feedback

Jonas Gehring, Kunhao Zheng, Jade Copet, Vegard Mella, Quentin Carbonneaux
Taco Cohen, Gabriel Synnaeve

{gehring,gab}@meta.com

Overview

A comprehensive journey through the RLEF methodology

01 Introduction & Motivation

LLM agents, execution feedback challenges, and the need for iterative code synthesis

02 The RLEF Method

Iterative code synthesis, MDP formulation, PPO optimization, and hybrid value function approach

03 Experimental Setup

CodeContests benchmark, Llama 3 models, evaluation metrics, and training configuration

04 Main Results

State-of-the-art performance, sample efficiency gains, and competitive comparisons

05 Inference-Time Behavior

Single vs multi-turn analysis, error recovery, code changes, and random feedback ablations

06 Ablation Studies

Few-shot vs SFT vs RLEF, single-turn training, repair models, and design validation

07 Related Work

Code generation LLMs, agentic frameworks, and prior RL approaches

08 Conclusions & Impact

Key contributions, limitations, broader impact, and reproducibility

LLM Agents and Execution Feedback

Two Crucial Skills for LLM Agents

1. Intent Deduction

Accurately understanding user intent when prompted, achieved through instruction fine-tuning according to user preferences (Ouyang et al., 2022; Rafailov et al., 2023)

2. Feedback Grounding

Taking into account feedback on intermediate actions to arrive at desired outcomes. **Crucial for grounding generations in concrete situations.**

Key Insight

Frame code generation as an **iterative task** with repeated attempts and execution feedback

→ Actions = code solutions

→ Observations = execution feedback

→ Reward = binary (tests pass/fail)

→ Optimization = end-to-end RL

</> The Code Generation Challenge

Problem: SOTA LLMs struggle to iteratively improve code compared to independent sampling (Kapoor et al., 2024; Xia et al., 2024)

Feedback Sources: Error messages, unit test results, compilation errors, runtime performance

Gap: Utilizing execution feedback has failed to yield substantial improvements when considering computational demands

Autonomous LLM Applications

- Web interaction:** Query search engines, interact with websites (Yao et al., 2022)

- Information retrieval:** Ensure accurate answers with up-to-date data (Mialon et al., 2024)

- Software development:** Generate code from high-level descriptions (Yang et al., 2024)

Complete Framework Overview



Two Test Sets

Public Tests

Provide execution feedback during attempts; accessible during iterative generation

Private Tests

Determine final correctness; compute binary reward for policy optimization

Key Benefits

- **Guards against shortcuts** — prevents copying expected outputs
- **Accelerates iteration** — limited public tests speed up generation
- **Enables RL training** — provides reward signal for optimization
- **Single model** — both synthesis and repair capabilities

Conversation Flow

Problem → Generate → Execute → Feedback → Improve → Repeat → Submit → Reward → Update

Turn Limits

Episode terminates when public tests pass OR specified turn limit reached (default: 3 attempts)

Multi-Turn Conversation Setup

Conversation Flow

1 Initial Prompt

Start dialog with problem description, query LLM for initial solution

2 Execute & Verify

Verify solution against public test set → passed/failed tests, errors

3 Format Feedback

If any public test fails, format feedback and append to dialog

4 Query Update

Query LLM for updated solution with full context (problem + prev solutions + feedback)

5 Submit Final

What additional context or dialog state do you need for this step?

Two Test Sets: Why Separate?

1. Guard Against Shortcuts

If test inputs/outputs are fixed, held-out tests prevent optimization shortcuts where LLM copies expected outputs based on execution feedback

2. Accelerate Iteration

Running full test suite may be computationally demanding. Limited public tests speed up iterative generation while maintaining effectiveness

Dialog Context Structure

Turn 1

User: Problem description

Assistant: Code solution

Turn 2

User: Feedback + "Give it another try"

Assistant: Updated code

... repeat

Execution Feedback Types

✓ Wrong Answer

✓ Timeout

✓ Exception

✓ Out of Memory

RLEF: MDP Formulation & PPO

MDP Components

Policy π : Language model

Observations o_t : Problem description

Actions a_t : Textual responses (code)

Observations o_t : Past observations, actions, execution feed back

Termination: Public tests pass OR turn limit reached

Reward: Binary (all tests pass/fail)

Reward Function

$$R(s_t, a_t) = r(s_t, a_t) - \beta \log[\pi(a_t|c_t)/\rho(a_t|c_t)]$$

$r(s_t, a_t) = +1$ (all tests pass)

$r(s_t, a_t) = -1$ (any test fails)

$r(s_t, a_t) = -0.2$ (invalid code)

$\beta = 0.05$ (KL regularization)

Hybrid Architecture

Token-Level Policy

Model policy at token level for fine-grained optimization

Turn-Level Value

Predict response value from last token of prompt; single advantage per response

Early experiments show this hybrid approach works best

PPO Optimization

Algorithm: Proximal Policy Optimization

Baseline: Value function $V(c_t)$

Advantage: $A_t = -V(c_t) + \sum R(s_i, a_i)$

Learning rate: $2e-7$ (AdamW)

Discount γ : 1 (no discounting)

KL penalty: Geometric mean (not product)

Training updates: 12k (8B), 8k (70B)

GPUs: 288 (8B), 2,304 (70B)

Failure Mode Addressed

Issue: Invalid code in non-final responses

Solution: Small penalty (-0.2) for invalid responses

KL Penalty Design

Geometric mean of token probabilities counters bias toward shorter generations, especially for non-final responses

CodeContests Benchmark & Models

CodeContests Benchmark

Competitive programming problems with natural language descriptions, public/private tests. High difficulty: algorithms, data structures, runtime efficiency.

13,328

Training

117

Validation

165

Test

Test Set Characteristics

Public Tests

1–7 test cases (median: 1) for feedback

Private Tests

Hidden tests for final correctness evaluation

Llama 3 Models

Strong instruction-following and code generation capabilities.

Llama 3.0 Instruct

8B, 70B parameters | Baseline models

Llama 3.1 Instruct

8B, 70B parameters | Primary models with enhanced coding

Training Configuration

Language: Python 3

Turn Limit: 3 attempts

Updates: 12k (8B), 8k (70B)


Time: ~20 hours

Evaluation Metric

n@k Average Solve Rate

Expectation that any of **n** solutions from **k** samples is correct. Enables fair comparison across sample budgets.

State-of-the-Art Performance



Key Achievements

70B Model: New SOTA

Beats AlphaCodium+GPT-4 (38.0 vs 29) with 1 vs 100 samples!

8B Model: Competitive

Matches AlphaCode 9B (16.0 vs 13.3) with 3 vs 1,000 samples!

10x Sample Reduction

Achieves SOTA with an order of magnitude fewer samples

Notable Comparisons

- **vs GPT-4:** RLEF 70B (37.5) with 3 samples vs GPT-4 Agent (29) with 100 samples
- **vs AlphaCode 2:** RLEF 70B (37.5) vs estimated AlphaCode 2 (34.2) on valid set
- **100 samples:** RLEF 70B reaches 54.5 (valid) and 54.5 (test) — beats all prior work

+94%

8B Valid
8.9 → 17.2

+45%

8B Test
10.5 → 16.0

+45%


70B Valid
25.9 → 37.5

+46%

70B Test
27.5 → 40.1

Single vs Multi-Turn Analysis

1@3 Solve Rates: Single vs Multi-Turn



Key Findings

Base Models: Rarely benefit from multi-turn feedback; sometimes perform worse

RLEF Models: Effectively leverage feedback to achieve larger gains

GPT-4o: Also shows stronger performance with independent sampling

Generalization

- **HumanEval+:** Improvements carry over with different feedback formatting
- **MBPP+:** Notable improvements despite simpler programming questions

8B CodeContests

Base ST: 11.8% | MT: 9.7%

RLEF ST: 16.0% | MT: 27.4%

Gain: +17.7% from multi-turn

70B CodeContests

Base ST: 25.3% | MT: 10.5%

RLEF ST: 40.1% | MT: 27.4%

Gain: +17.1% from multi-turn

Cross-Domain


HumanEval+: 80.7% (+3.8%)

MBPP+: 71.7% (+3.2%)

Generalizes to new domains

Error Recovery & Code Changes

Error Analysis Over 20 Rollouts



8B model results (70B similar)

Random Feedback Ablation

With random feedback, self-repairs are severely impaired, proving RLEF models meaningfully leverage true execution feedback.

17.2

True Feedback

12.2

Random Feedback

Key Observations

1. Fewer Initial Errors

RLEF-trained models produce fewer wrong outputs in first response but are more prone to timeout (efficiency focus)

2. Better Error Recovery

Significantly improved recovery from all error categories (output, exception, timeout, OOM) in subsequent responses

3. Larger Code Edits

Higher chrF scores indicate more substantial changes. Base models repeat same code despite feedback

Pass@1 vs Pass@10 Analysis


Random feedback drops pass@1 significantly (less targeted repair) but affects pass@10 less (can still sample diverse solutions)

→ **Pass@1:** Measures precision of arriving at correct solution

→ **Pass@10:** Measures recall (whether any solution passes)

Learning Iterative Code Synthesis

Methods for Iterative Capabilities



1@3 solve rates on CodeContests (Instruct models)

Few-Shot Prompting

Finding: Detrimental to Instruct models. Base models also achieve lower performance than zero-shot.

Supervised Fine-Tuning (SFT)

Filter 313k successful trajectories from Llama 3.1 70B → fine-tune Base & Instruct models.

10.3

8B SFT Valid

17.2

8B RLEF Valid

Result: Moderate improvements on validation only. RLEF significantly outperforms SFT.

Single-Turn vs Multi-Turn Training

Compare traditional single-generation vs iterative setup with same training loop.

8B Model

Single-turn hurts performance on test set

70B Model

Benefits from single-turn training; shows transfer to multi-turn inference

Positioning RLEF in Code Generation

Agentic Frameworks

AlphaCodium, MapCoder, Reflexion, LDB: Rich manual scaffolding, chaining several LLM calls with code execution.

Issues:

- **Inference cost:** Dozens of LLM calls per solution
- **Independent sampling competitive** (Kapoor et al., 2024)
- **Large models required** for effective feedback

Prior RL Work

Le et al. (2022), Xu et al. (2024): Train code LLMs with scalar rewards from unit tests.

- **CodeRL:** Policy gradients + next-token loss on rewards
- **DeepSeek-Coder:** Binary reward from unit tests

Limitation: No textual feedback during generation

RLEF Key Distinctions

1. Textual Execution Feedback

Not just scalar rewards — provides rich textual feedback for code synthesis and repair in a single model

2. End-to-End RL Training

Optimizes for leveraging feedback through multi-turn conversation, not just single generation

3. Sample Efficiency Focus

Shifts from large-sample inference to high accuracy with low sample budgets

Concurrent Work: SCoRe

Kumar et al. (2024): Two-stage RL for self-correction; outputs two successive solutions.

Key Difference:

SCoRe doesn't leverage execution feedback at inference time, limiting its ability to adapt to new environments.

Impact, Limitations & Future

Key Contributions

1. RLEF Method

End-to-end RL for LLMs to leverage execution feedback, grounding future generations in environment feedback

2. SOTA Performance

Substantial improvements on CodeContests; reduces sample budget by 10x

3. Generalization

Improved multi-turn performance carries over to HumanEval+ and MBPP+

Performance Breakthrough

54.5%

70B 10@100

40.1%

70B 1@3

Limitations

1. Single Solution Focus

Limited to improving single solution; doesn't handle task decomposition

2. Requires Test Cases

Iteration requires test cases, which may not be readily available

3. Domain Specificity

Training on competitive programming may limit broader applicability

Broader Impact

Positive

Amplifies utility for software development and quality control

Consideration

Requires quality control and guard-railing to promote safety

Safety

Execution confined to local sandboxes; follows AI agent governance frameworks

Future Directions

Combine with automatic unit test generation; extend to larger tasks requiring decomposition

RLEF Enables Effective Feedback Grounding



RLEF opens new possibilities for autonomous code generation and repair, shifting focus from large-sample inference to high