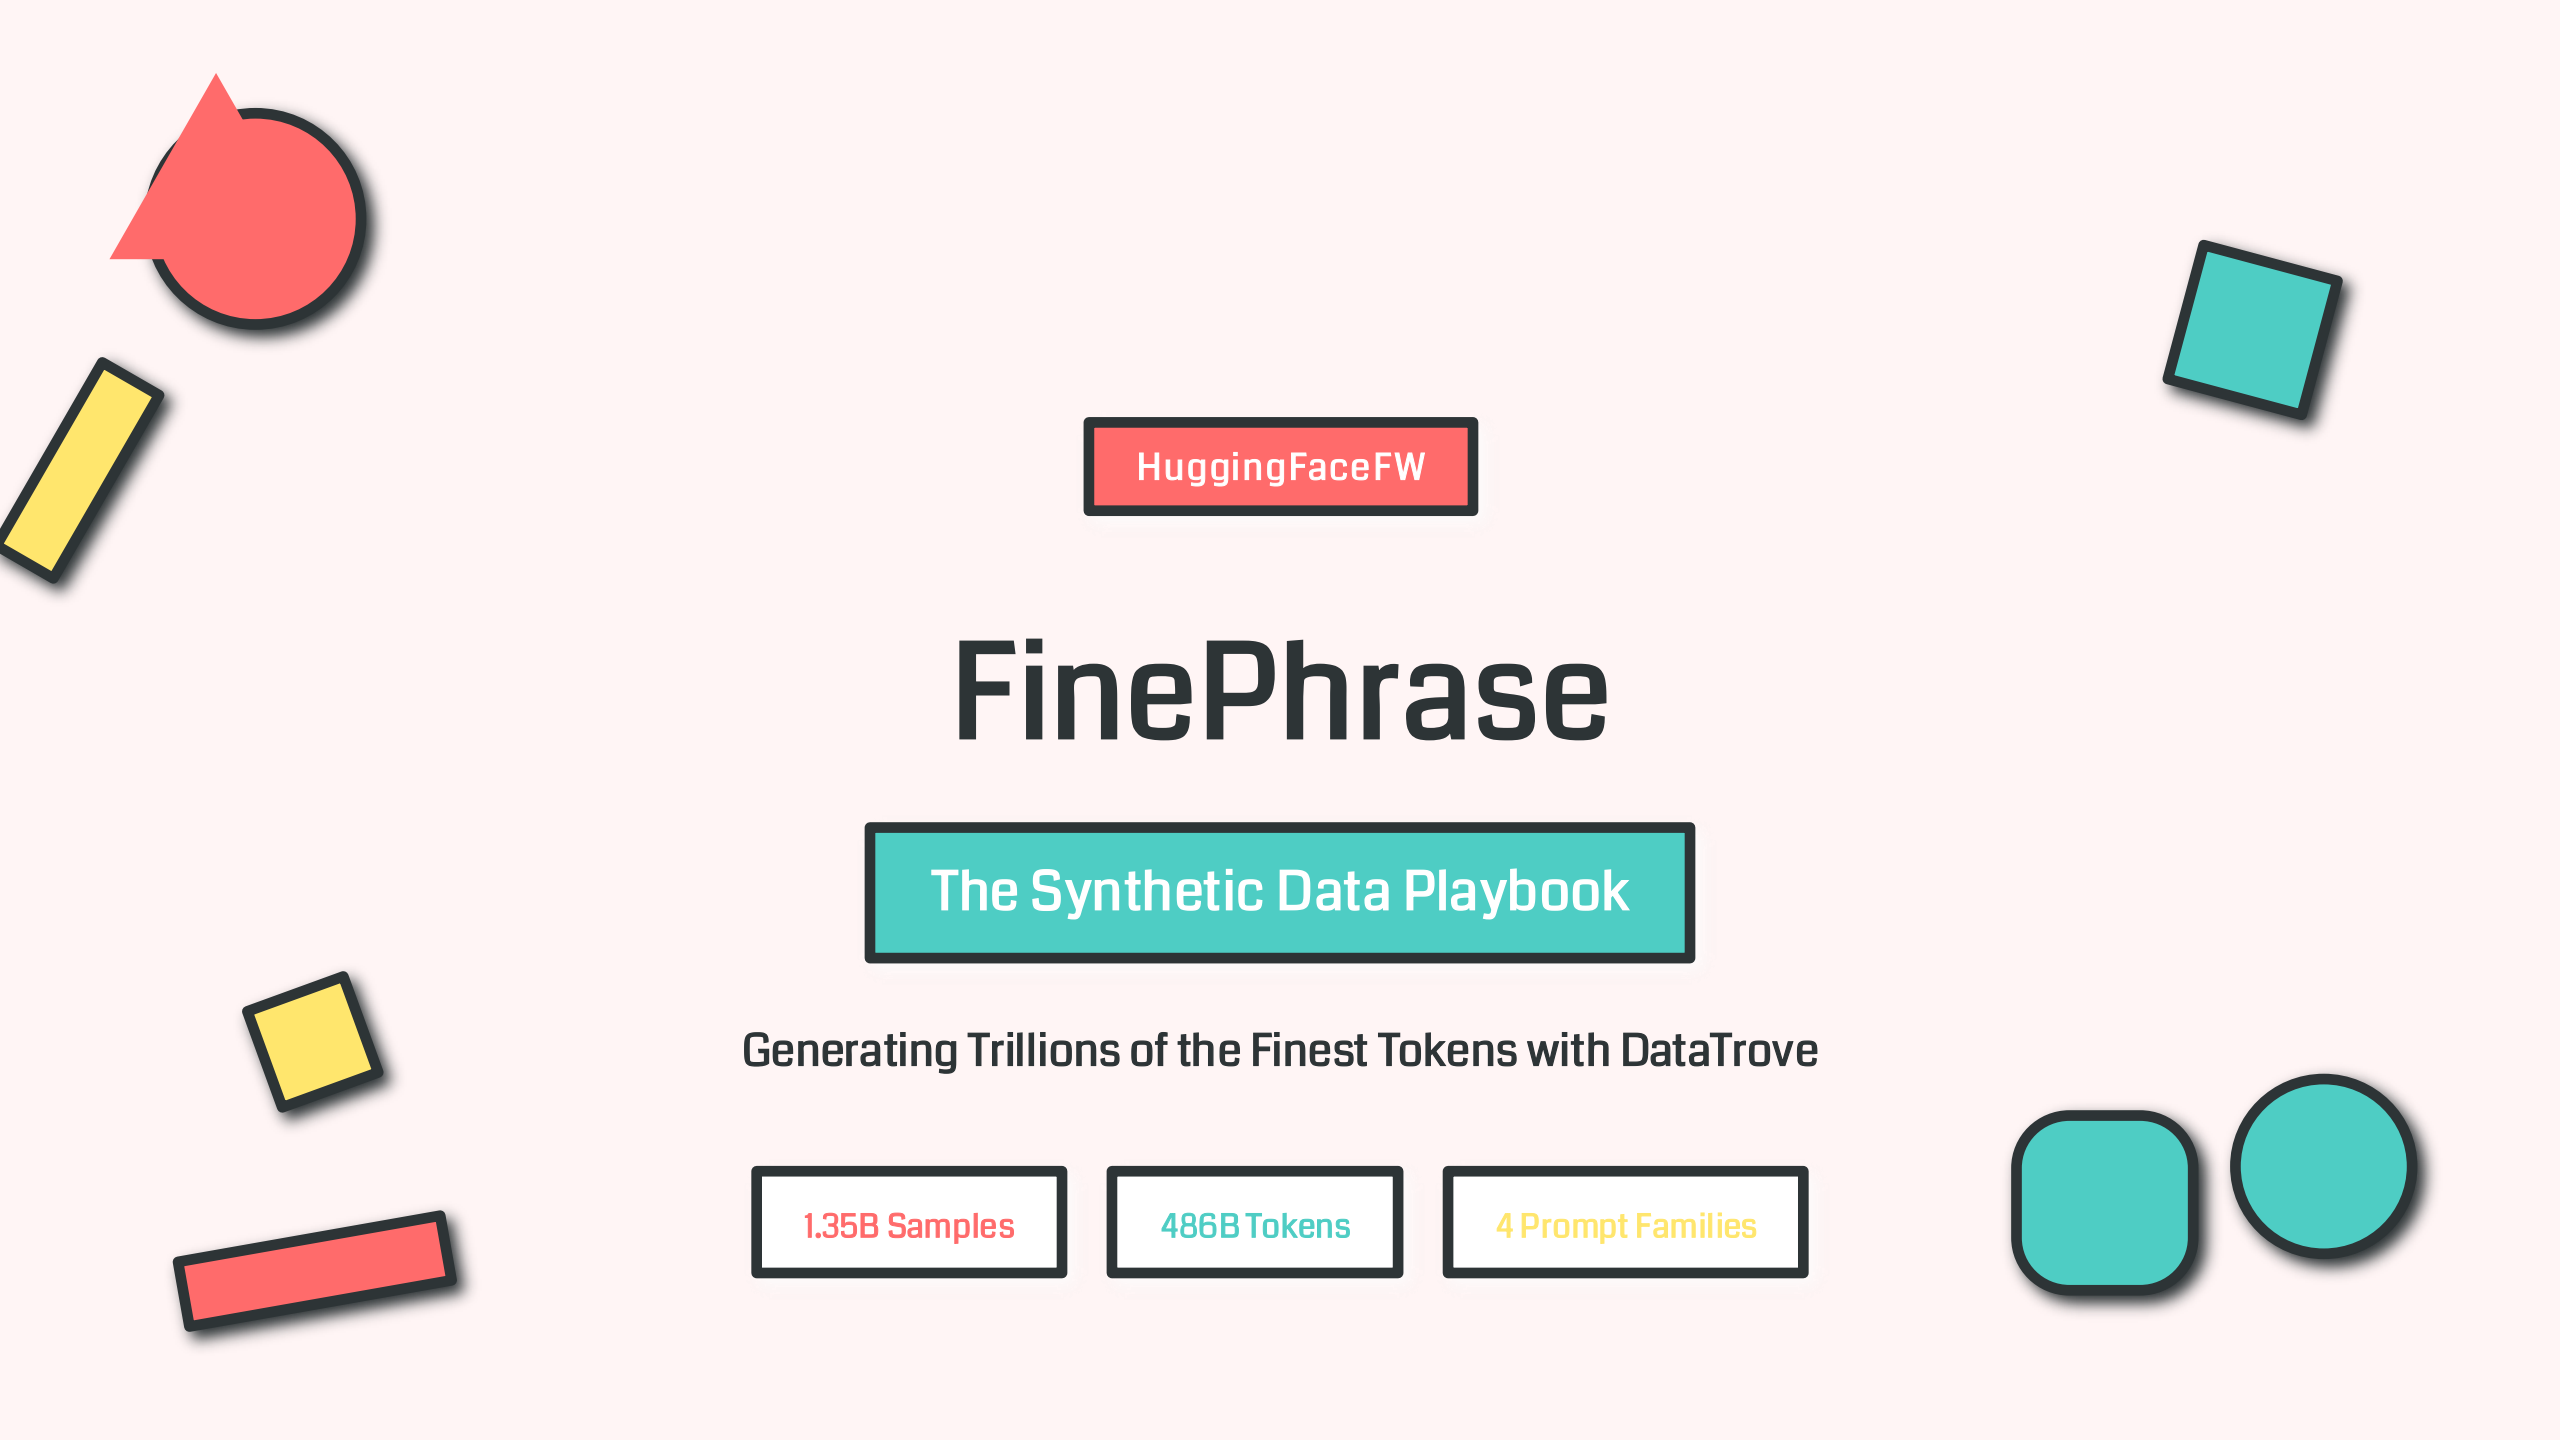




HuggingFaceFW

FinePhrase

The Synthetic Data Playbook


Generating Trillions of the Finest Tokens with DataTrove

1.35B Samples

486B Tokens

4 Prompt Families

Contents



01

Project Overview

What is FinePhrase & its purpose

03

Generation Model

SmolLM2-1.7B-Instruct configuration

05

Data Processing Pipeline

From input to synthetic output

07

Data Schema & Structure

Document format & metadata fields

09

Usage & Access

How to load & use FinePhrase

11

Limitations

Known issues & considerations

02

Technical Architecture

DataTrove pipeline & infrastructure

04

Prompt Family Design

FAQ, Math, Table, Tutorial prompts

06

Scale & Statistics

Generation metrics & token counts

08

Quality Assurance

Validation & filtering mechanisms

10


Research Impact

Applications & best practices

12

Conclusion

Future of synthetic data



01

Project Overview

What is FinePhrase?



Massive Synthetic Dataset

FinePhrase transforms educational web content from **FineWeb-Edu** into structured formats using AI. It represents one of the largest synthetic data generation projects for LLM training.



Four Transformation Types

Content is transformed into FAQ, Math problems, Structured Tables, and Tutorials—each serving different training objectives for language models.



Powered by DataTrove

Built on Hugging Face's **DataTrove** pipeline library, enabling scalable, platform-agnostic data processing from local machines to SLURM clusters.

Generation Scale

Source Documents

339.3M

Output Samples

1.35B


Completion Tokens

486.4B

4x Expansion Rate

The FineData Ecosystem

FinePhrase within the HuggingFaceFW Family



FineData Team

Hugging Face Science Team



FineWeb

15T tokens English dataset for LLM pre-training

15 Trillion



FineWeb-Edu

Filtered educational content from FineWeb

3.5 Billion



FineWeb2

Extension to 1000+ languages

1000+ Langs



FinePDFs

3T tokens from PDF documents

3 Trillion



FineWiki

Updated Wikipedia in 300+ languages

300+ Langs



FinePhrase

Synthetic transformation of educational content

486B Tokens

FineWeb-Edu → FinePhrase



Synthetic Augmentation

DataTrove: The Processing Engine

Platform-Agnostic Pipeline Processing Library



What is DataTrove?

DataTrove is a Python library for **processing, filtering, and deduplicating** text data at very large scale. It provides prebuilt processing blocks with a framework for custom functionality.

⚡ Low memory usage & multi-step design

Core Capabilities



Readers
Multiple formats



Filters
SOTA quality



Deduplication
MinHash, ExactSubstr



Writers
Flexible output



Tokenization
Multiple tokenizers



Inference
LLM generation

Execution Environments



LocalPipelineExecutor

Multi-CPU local execution



SlurmPipelineExecutor

SLURM cluster deployment



RayPipelineExecutor

Ray distributed compute

🌐 **Platform-agnostic: Same pipeline runs anywhere**

`</> pip install datatrove[inference]`

Generation Model

SmolLM2-1.7B-Instruct



Model Selection

HuggingFaceTB/SmolLM2-1.7B-Instruct

SmolLM2-1.7B-Instruct was chosen for its **optimal balance** of efficiency and capability. Despite its compact size, it delivers high-quality synthetic content across diverse transformation tasks.

Parameters

1.7B

Architecture
Transformer

Why This Model?



Efficient Inference

Smaller model enables faster generation at scale, reducing computational costs for 486B tokens.



Instruction-Tuned

Fine-tuned for following complex prompts—critical for FAQ, Math, Table, and Tutorial generation.



Quality Output

Produces coherent, structured content matching the educational quality of FineWeb-Edu source.

Model Capabilities



FAQ Generation

Extracts key questions and provides clear answers



Math Problems

Creates word problems with step-by-step solutions



Structured Tables

Organizes data into markdown tables



Generation Configuration

Parameters & Decoding Strategy

Core Parameters

Temperature

1.0

Creativity control

Top-p (Nucleus)

1.0

Full nucleus sampling

Top-k

50

Top-k sampling limit

Max Tokens

2048

Output length limit

Speculative Decoding

Method

Suffix

Speculative Tokens

32

Context & System Settings



Model Max Context

Total context window size

8192



System Prompt

Global instruction

None



Input Column

Source text field

text

Why These Settings?

- 1 Temp=1.0:** Balanced creativity—neither too random nor too deterministic
- 2 Top-p=1.0:** Full nucleus sampling for diverse outputs
- 3 Top-k=50:** Limits vocabulary to top 50 tokens per step
- 4 Speculative:** Speeds up generation by 2-3x

Source Dataset

FineWeb-Edu Sample-350BT



FineWeb-Edu

Educational Content Filter

FineWeb-Edu is a **filtered subset** of FineWeb containing the most educational content. It uses a classifier to identify high-quality educational documents from web data.

Full Dataset

3.5B

documents

Sample Used

350BT

config

Input Statistics



Dataset Name

HuggingFaceFW/fineweb-edu



Configuration

sample-350BT



Split

train

Why FineWeb-Edu?



High Quality

Classifier-filtered for educational value



Pre-cleaned

Already deduplicated and filtered



Diverse

Broad educational topics covered



Safe

PII and toxic content filtered

Prompt Family 1

FAQ Generation



FAQ Prompt Template

Rewrite the document as a comprehensive FAQ (Frequently Asked Questions). Extract or infer the key questions a reader would have about this topic, then provide clear, direct answers. Order questions logically, from foundational to advanced, or by topic area. Each answer should be self-contained and understandable without reference to other answers. Ensure the FAQ works as a standalone document. Output only the FAQ, no thing else.
Document: [[DOCUMENT]]

Key Design Elements

1

Extract & Infer

Both explicit and implicit questions

2

Logical Ordering

Foundational to advanced

3

Self-Contained

No cross-references needed

4

Standalone

Works independently

Training Objective



Question Answering

Teaches models to answer questions accurately



Information Extraction

Identifies key points in documents



Structured Output

Generates organized Q&A format

Generation Stats

Documents

338.97M

Tokens

148.1B

Mean Tokens/Sample

436.96

Prompt Family 2

Mathematical Word Problems



Math Prompt Template

Rewrite the document to create a mathematical word problem based on the numerical data or relationships in the text. Provide a step-by-step solution that shows the calculation process clearly. Create a problem that requires multi-step reasoning and basic arithmetic operations. It should include the question followed by a detailed solution showing each calculation step. Output only the problem and solution, nothing else.
Document: [[DOCUMENT]]

Key Design Elements

1

Numerical Focus

Extracts quantitative data

2

Step-by-Step

Detailed solution process

3

Multi-Step

Complex reasoning required

4

Basic Arithmetic

Accessible math operations

Training Objective



Mathematical Reasoning

Develops step-by-step problem solving



Numerical Understanding

Extracts numbers from text context



Chain-of-Thought

Shows explicit reasoning steps

Generation Stats

Documents

338.75M

Tokens

98.4B

Mean Tokens/Sample

290.51

Prompt Family 3

Structured Tables



Table Prompt Template

Rewrite the document as a structured table that organizes the key information, then generate one question-answer pair based on the table. First extract the main data points and organize them into a clear table format with appropriate headers using markdown table syntax with proper alignment. After the table, generate one insightful question that can be answered using the table data. Provide a clear, concise answer to the question based on the information in the table. Output only the table followed by the question-answer pair, nothing else.

Document: [[DOCUMENT]]

Key Design Elements

1

Data Extraction

Identifies key data points

2

Markdown Format

Standard table syntax

3

Proper Alignment

Structured headers

4

Q&A Pair

Tests table comprehension

Training Objective



Structured Data

Understands tabular formats



Information Retrieval

Extracts specific data points



Table QA

Answers based on table content

Generation Stats

Documents

338.55M

Tokens

92.4B

Mean Tokens/Sample

272.94

Prompt Family 4

Tutorial Generation



Tutorial Prompt Template

Rewrite the document as a clear, step-by-step tutorial or instructional guide. Use numbered steps or bullet points where appropriate to enhance clarity. Preserve all essential information while ensuring the style feels didactic and easy to follow. Output only the tutorial, nothing else.
Document: [[DOCUMENT]]

Key Design Elements

1

Step-by-Step
Numbered instructions

2

Bullet Points
Visual clarity aids

3

Didactic Style
Teaching-oriented tone

4

Easy to Follow
Clear progression

Training Objective



Instruction Following

Learns to follow procedural instructions



Sequential Reasoning

Understands ordered steps



Educational Content

Produces teaching materials

Generation Stats

Documents

337.78M

Tokens

147.4B

Mean Tokens/Sample

436.49

Generation Scale

The Numbers Behind FinePhrase

Overall Statistics



Source Documents

339.3M

Input count



Output Samples

1.35B

Generated total




Completion Tokens

486.4B

Total generated

4x Expansion Explained



Each source document is processed through **4 different prompt families**, creating 4 unique synthetic samples per input document.

Scale Context



Web Scale

Equivalent to processing **all Wikipedia articles** 50+ times



Compute Time

Generated using **distributed GPU clusters** over weeks



Storage


Hundreds of **terabytes** of parquet files

One of the largest synthetic datasets ever created

Configuration-Level Statistics

Detailed Breakdown by Prompt Family

Generation Metrics by Config



FAQ
436.96
mean tok

Math
290.51
mean tok

Table
272.94
mean tok

Tutorial
436.49
mean tok

 Overall Mean: 359.20 tokens/sample

Key Insights

Longest Output

FAQ & Tutorial produce the longest outputs (~437 tokens)

Most Concise


Table format is most compact (~273 tokens)

Balanced

Math strikes middle ground (~291 tokens)

Data Schema

Document Structure & Metadata Fields



Top-Level Fields



id
string

Unique identifier for each sample



text
string

Source input from FineWeb-Edu (NOT generated output)



rollout_results
list

List of generation result objects (one per rollout)

rollout_results Structure



finish_reason

Why generation stopped (e.g., "stop", "length")



text

Generated output (e.g., FAQ, Math problem)



usage

- completion_tokens
- prompt_tokens
- prompt_tokens_details
- total_tokens




Each sample preserves the original source text alongside generated outputs

Access generated content via: rollout_results
[0].text

Data Processing Pipeline

From Input to Synthetic Output

Pipeline Flow



Pipeline Block Types

-  **Readers**
Read data from different formats (JSON, CSV, Parquet)
-  **InferenceRunner**
Handles LLM generation with async batching
-  **Filters**
Remove documents based on quality criteria
-  **Writers**
Save to disk/cloud in various formats

InferenceRunner Features

-  **Async Batching**
Maintains high GPU utilization
-  **Checkpointing**
Resumable generation for fault tolerance
-  **Multi-Node**
Distributed across SLURM clusters
-  **Server Support**
vLLM, SGLang, OpenAI-compatible

Inference Configuration

Execution & Batching Strategy

Supported Server Types



vLLM

High-throughput inference



SGLang

Structured generation



OpenAI-Compatible

API endpoints



Dummy

Testing & debugging

Key Configuration Parameters



rollouts_per_document

Multiple generations per sample

1



max_concurrent_generations

Batching control for throughput

500



checkpoints_local_dir

Resumable checkpoint location

/fsx/...

Async Batching Benefits



High GPU Utilization

Keeps GPUs busy with continuous requests



Reduced Latency

Overlaps preprocessing with generation



Scalability

Handles billions of samples efficiently



Configurable for 1B to 1T parameter models

Quality Assurance

Validation & Filtering Mechanisms

Quality Control Measures



finish_reason Tracking

Monitors why generation stopped: "stop" = success, "length" = truncated



Invalid Request Filtering

Per-config totals slightly below source count due to **skipped invalid requests**



Context Budget Management

Some long inputs **truncated** to satisfy context budgets (8192 tokens)

⚖️ Balance between throughput and quality maintained throughout generation

Quality Metrics

Success Rate

~99.9%




Truncation Rate

<1%



Invalid Skip Rate


<0.5%



🛡️ Robust error handling ensures reliable generation

Data Retention & Processing

Policies for Reproducibility & Efficiency



Retention Strategy



Source Documents Retained

Original FineWeb-Edu documents **preserved for reproducibility** and provenance tracking



Generated Outputs Stored

All synthetic outputs saved with **full metadata** including generation parameters



Complete Metadata

Token counts, finish reasons, and usage statistics **preserved for analysis**

Fault Tolerance

✓ skip_completed

Skip already completed tasks on restart

🔄 Efficient Reprocessing

Only failed tasks rerun on relaunch

Completion Tracking

📁 Marker Files

Empty files in logging_dir/completions folder

✅ Task-Level Tracking

Each task writes completion marker when done

▶ Resume Capability

Relaunch same executor to continue

∞ Designed for long-running distributed jobs

Anonymization & Privacy

Considerations for Synthetic Data

Privacy Protection Measures



Pre-Generation Filtering

FineWeb-Edu source data undergoes **PII filtering** before FinePhrase generation



Synthetic Transformation

Content transformation helps **reduce privacy risks** by restructuring information



Educational Focus

Source data focuses on **educational content**, minimizing sensitive information

Related: FinePDFs uses PII anonymization (emails, IPs)

Known Limitations



Hallucination Risk

Outputs may contain fabricated information



Pattern Reproduction

May inadvertently reproduce training patterns



No Guaranteed Anonymity

Synthetic $\neq$ anonymous; review before use



Balance utility and privacy carefully

Usage & Access

How to Load FinePhrase



Loading FinePhrase

```
from datasets import load_dataset
# Load all configurations
ds_all = load_dataset("HuggingFaceFW/finephrase", "all")
# Load specific configurations
ds_faq = load_dataset("HuggingFaceFW/finephrase", "faq")
ds_math = load_dataset("HuggingFaceFW/finephrase", "math")
ds_table = load_dataset("HuggingFaceFW/finephrase", "table")
ds_tutorial = load_dataset("HuggingFaceFW/finephrase", "tutorial")
```

Available Configurations

all

Combined

faq

Q&A

math

Problems

table

Tables

tutorial

Guides

Access Methods

**Hugging Face Datasets**

pip install datasets

**Hugging Face Hub**

huggingface_hub library

**DataTrove**

pip install datatrove

Streaming supported for large-scale processing

Reproducibility & Open Science

Transparent Research Practices

Available Resources



Main Scripts

`examples/inference/finephrase.py` and `generate_data.py`



Full Configuration

Complete generation parameters **documented and versioned**



Transparent Pipeline

DataTrove pipeline is **open-source** and auditable

Licensing

ODC-By License

Open Data Commons Attribution License

Free for research & development

Hugging Face Commitment



Open Data

Releasing large-scale datasets for LLM development



Community

Accelerating open LLM research worldwide



Science Team

Part of Hugging Face Science initiative



Democratizing AI through open data

Research Applications & Impact

Use Cases & Best Practices

Research Use Cases

Pre-training

Augment base model training data

Math Reasoning

Improve step-by-step problem solving

Tutorial Following

Learn procedural instruction

Instruction Tuning

Train QA and reasoning capabilities

Structured Data

Enhance table understanding

Data Augmentation

Expand training set diversity

 Synthetic data complements real data in LLM training

Best Practices

Mix with Real Data

Combine synthetic and authentic sources

Task-Specific

Use for targeted augmentation

Validate Outputs

Review before deployment

Monitor Quality

Track for degradation over time

 Based on synthetic data literature best practices

Limitations & Known Issues

Important Considerations

Documented Limitations



Hallucinations

Outputs are **model-generated** and may contain fabricated information



Input Truncation

Some long inputs **truncated** to satisfy 8192 token context budget



Skipped Requests

Per-config totals slightly below source due to **invalid request skipping**

Additional Concerns



Quality Variance

Output quality varies across prompt families



Feedback Loops

Risk of model collapse if used iteratively

Mitigation Strategies



Validate Outputs

Review synthetic data before training



Mix Data Sources

Combine with real data to prevent collapse



Monitor Quality


Track model performance on validation sets



Transparent documentation of limitations

Best Practices

Recommendations for Using FinePhrase



1 Mix Synthetic with Real

Always combine synthetic data with **real data** to prevent model collapse and maintain diversity.

💡 Recommended ratio: 10-25% synthetic

2 Task-Specific Usage

Use synthetic data for **specific task augmentation** rather than sole training source.

🎯 Target: QA, reasoning, structured data

3 Validate Before Deployment

Review outputs for quality and accuracy before using in production models.

✓ Sample and inspect generated content

4 Monitor & Track Provenance

Track data sources in training mixtures to identify potential circular dependencies.

📝 Log data composition in experiments

★ Following best practices ensures safe and effective use of synthetic data

Technical Implementation

Implementation Details & Infrastructure

Token Counting Methodology



Counting Script

`examples/inference/count_completion_tokens.py`



Stats Location

`projects/datatrove/finephrase_token_counts/`

Infrastructure & Optimization



Parquet Format

Efficient columnar storage



SLURM Clusters

Distributed processing



Memory Optimization

Low memory footprint



Async Processing

High-throughput generation

Scale Metrics

Processing Time

Distributed GPU cluster time

Weeks

Storage

Parquet file storage

100s TB

Parallel Tasks

Concurrent generation tasks

1000s



Built for billion-scale generation

Citation & Attribution

How to Cite FinePhrase



BibTeX Citation

```
@misc{niklaus2026_the_synthetic_data_playbook_generating_trillions_of_the_finet_tokens,  
title={The Synthetic Data Playbook: Generating Trillions of the Finest Tokens},  
author={  
Joel Niklaus and Guilherme Penedo and  
Hynek Kydlíček and Elie Bakouch and  
Lewis Tunstall and Ed Beeching and  
Thibaud Frère and Colin Raffel and  
Leandro von Werra and Thomas Wolf  
},  
year={2026},  
}
```

Authors

Joel Niklaus

Guilherme Penedo

Hynek Kydlíček

Elie Bakouch

Lewis Tunstall

Ed Beeching

Thibaud Frère

Colin Raffel

Leandro von Werra

Thomas Wolf

Licensing Information



ODC-By

Open Data Commons Attribution

Free to use for research and development with attribution requirements.



Permitted Uses

- ✓ Research
- ✓ Commercial development
- ✓ Model training
- ✓ Distribution



Please cite when using in research

The Future of Synthetic Data

FinePhrase demonstrates the **power of systematic synthetic data generation** at scale. As LLMs continue to evolve, synthetic data pipelines like this will be crucial for creating diverse, high-quality training data.

The combination of **DataTrove's scalable infrastructure** with **careful prompt engineering** opens new possibilities for data augmentation in AI development.



1.35B Samples

Massive scale generation



DataTrove

Scalable infrastructure



4 Prompt Families

Diverse transformations



Empowering the Next Generation of AI