

Nemotron 3 Super


Open, Efficient Mixture-of-Experts Hybrid Mamba-Transformer Model for Agentic Reasoning

120B Total Params

12B Active


1M Context

The first model pre-trained in NVFP4 with LatentMoE architecture and native speculative decoding





Presentation Overview



01

Executive Summary

Key innovations, specifications, and open source release

02

Model Architecture

LatentMoE, MTP, and hybrid Mamba-Attention design

03

Pretraining Methodology

NVFP4 training, data curation, and optimization

04

Post-Training Pipeline

SFT, RLVR, SWE-RL, and RLHF stages

05

Quantization & Inference

FP8 and NVFP4 deployment strategies

06


Performance Evaluation

Comprehensive benchmark results and comparisons

07

Conclusion & Open Source Impact



Summary of contributions and community resources




CHAPTER 01

Executive Summary

Key innovations and performance highlights of Nemotron 3 Super - the next generation of efficient, agentic AI





Nemotron 3 Super at a Glance

120B

Total Parameters

Sparse MoE Architecture

12B

Active Parameters

Per Forward Pass

1M

Context Length

Maximum Tokens

Key Innovations

- **First NVFP4 Pretraining:** Stable training at 25T tokens in 4-bit precision
- **LatentMoE Architecture:** Optimizes accuracy per FLOP and per parameter
- **MTP Layers:** Native speculative decoding for inference acceleration
- **Hybrid Design:** Mamba-2 + Attention with strategic global anchors

Inference Throughput Gains

2.2× vs GPT-OSS-120B
8k input / 64k output


7.5× vs Qwen3.5-122B
Same sequence lengths



Training Scale: 25 Trillion Tokens | 88 Layers | 512 Experts | Top-22 Routing



Three Key Innovations



1

LatentMoE Architecture

Hardware-software co-design for optimal accuracy per FLOP and per parameter. Projects tokens into lower-dimensional latent space ($\ell < d$) for routing and expert computation, reducing memory bandwidth and communication by factor d/ℓ . Savings enable scaling both total experts ($N' = N \cdot d/\ell$) and active experts ($K' = K \cdot d/\ell$), achieving higher quality at similar inference cost.

2

NVFP4 Pretraining

First large-scale model trained in 4-bit floating point from scratch. Uses E2M1 element format with 16-element micro-blocks and E4M3 micro-block scaling. 2D block scaling for weights ensures consistency between forward and backward passes. Demonstrates stable training at 25T tokens with strategic BF16/MXFP8 layers for critical components.

3

Multi-Token Prediction (MTP)

Shared-weight design enables robust autoregressive drafting for native speculative decoding. Unlike independent offset heads, shared parameters regularize across prediction horizons, improving stability with self-generated hidden states. Achieves **3.45 average acceptance length** on SPEED-Bench (draft length 7), outperforming DeepSeek-R1 and competitive with Qwen3-Next.



Open Source Release

Model Checkpoints

Nemotron 3 Super 120B-A12B NVFP4

Post-trained & NVFP4 quantized

Nemotron 3 Super 120B-A12B FP8

Post-trained & FP8 quantized

Nemotron 3 Super 120B-A12B BF16

Post-trained model

Nemotron 3 Super 120B-A12B Base BF16

Base model (pre-SFT)

Qwen3-Nemotron-235B-A22B-GenRM-2603

GenRM used for RLHF

Pretraining Datasets

Nemotron-Pretraining-Specialized-v1.1

- Synthetic code concepts & algorithms
- Formal logic problems & solutions
- Economics multiple-choice questions
- Multiple choice question augmentation

Post-Training Datasets



Nemotron-Super-Post-Training-Data

- RL environments (21 types)
- SFT datasets for agentic capabilities
- Multi-turn conversation data



Training Recipe: github.com/NVIDIA-NeMo/Nemotron



n



CHAPTER 02

Model Architecture & Innovations

Deep dive into LatentMoE, Multi-Token Prediction, and the hybrid Mamba-Attention design





Architecture Overview

88

Total Layers

4096

Model Dim

512

Total Experts

22

Top-K Routing

Attention Configuration

Query Heads (n_q)

32

KV Heads (n_kv)

2

Head Dimension

128

GQA: Grouped-Query Attention for efficient KV cache

Mamba Configuration

State Dimension

128

Mamba Groups

8

Mamba Heads

128

Head Dimension

64

MoE Configuration

Expert Hidden Dim

2688

Shared Expert Size

5376

MoE Latent Size

1024

MTP Configuration

MTP Layers

2

Weight Sharing

Yes

Enables native speculative decoding



Hybrid Pattern: Mamba-2 + LatentMoE + Attention layers with strategic global anchors



LatentMoE: Hardware-Aware Expert Design

The Problem with Standard MoE

Existing MoE designs focus on **accuracy per FLOP** but neglect critical deployment constraints:

- ▶ **Memory bandwidth** bottlenecks in low-latency serving
- ▶ **All-to-all communication** overhead in distributed inference
- ▶ **Sharding overhead** from large hidden dimensions

Design Principles

1. Memory Bandwidth Dominates

In low-latency serving, reduce d or m

2. All-to-All Dominates

In throughput serving, reduce d or K

3. Preserve Nonlinear Budget

Keep $K \cdot m$ fixed for quality

Additional Principles

4. Feature Rank Lower Bound

Task-specific effective feature rank r_{eff} imposes minimum on d reduction

5. Scale N and K Exponentially

Scaling both total experts N and top- K per token exponentially expands expert combination space


Key Insight

Hidden dimension d is the most promising axis for reduction. Increase K by factor α and reduce d by same factor $\alpha \rightarrow$ higher accuracy at similar inference cost.




LatentMoE Architecture Detail

Standard MoE



LatentMoE



Memory Savings

Per-expert weight loads reduced by factor d/ℓ

Communication Savings

All-to-all payloads reduced by factor d/ℓ

Quality Gains

Scale N' and K' by $d/\ell \rightarrow$ higher accuracy

✓ Non-routed computations (router, shared expert, non-expert layers) remain in full dimension d



Multi-Token Prediction (MTP)

Training Benefits

Multi-Step Dependencies

Optimizes model to predict multiple future tokens at each position

Longer-Range Structure

Encourages representations capturing extended dependencies

Quality Improvements

Consistent gains in validation loss and downstream accuracy

Inference Benefits

Native Speculative Decoding

Auxiliary heads function as internal draft model

Single Forward Pass

Candidates verified in one pass with minimal additional FLOPs

Reduced Latency

Substantially reduces decoding latency vs external draft model

Shared-Weight Design for Robust Autoregressive Drafting

Problem with Standard MTP:

- Independent heads trained for fixed offsets
- Limits speculative decoding to N draft tokens
- Training-inference mismatch with self-generated states
- Acceptance rates degrade with draft length

Nemotron 3 Super Solution:

- Share parameters across MTP heads
- Unified prediction head exposed to multiple offsets
- Regularizes across prediction horizons
- Same head applied recursively for longer drafts
- Milder acceptance degradation at depth



MTP Performance Results



SPEED-Bench Results (Draft Length = 7)

Nemotron 3 Super	3.45
Qwen3-Next	3.33
DeepSeek-R1	2.70
Highest overall average acceptance length across all domains	

Acceptance Rate by Draft Index

Index 1-3	N3S > Qwen3 ≈ DS-R1
Index 4-5	N3S > Qwen3 > DS-R1
Index 6-7	N3S >> DS-R1 (large gap)
Pronounced gap at indices 4-7 where recursive drafting is most challenging	

Coding Domain

3.78
Nemotron 3 Super

Math Domain

3.73
Nemotron 3 Super

Reasoning Domain

3.78
Nemotron 3 Super



MTP shifts throughput-latency Pareto frontier: D=1 to D=3 draft depths deliver higher aggregate TPS for any given median user latency



Hybrid Interleaved Architecture

The Challenge

Quadratic KV Cache Growth

Self-attention layers have $O(n^2)$ memory complexity, creating severe bottlenecks for long-context inference.

The Solution

Mamba-2 Blocks: $O(n)$ complexity

Strategic Attention Anchors

Sparse MoE for Capacity

Layer Pattern (88 Total Layers)



Periodic interleaving: MoE layers paired with Mamba-2 blocks, attention as global anchors

Mamba-2 Benefits


- Constant-sized state
- Linear-time generation
- Reduced memory overhead
- Lower latency

Attention Anchors

- Full-token interaction
- Long-range routing
- Global dependency modeling
- Strategic placement

Sparse MoE



- Top-22 routing
- 512 total experts
- 12.7B active params
- 120.6B total params



CHAPTER 03

Pretraining Methodology

Training approach, data curation, and optimization strategies at scale





NVFP4 Pretraining

NVFP4 Format

Element Format: E2M1

2-bit exponent, 1-bit mantissa

Micro-Blocks: 16 elements

E4M3 micro-block scaling

Global Scale: FP32

Second-level per-tensor

Quantization Strategy

Weights: 2D Block Scaling

Consistency between fprop & bprop

Activations/Gradients: 1D Blocks

Along GEMM reduction axis

RHT + Stochastic Rounding

On wgrad inputs & gradients

Layer-Specific Precision

NVFP4 (Default)

All linear layers unless noted

MXFP8

Mamba output projection (mitigates underflows)

BF16

Final 15%, latent projections, MTP, QKV, embeddings

FP32

Router, output layers



First model to demonstrate stable large-scale training in NVFP4 up to 25T tokens



Training Stability Analysis

Observed Phenomenon

Zero-Valued Gradients

7% of total parameters by end of pretraining

Magnitude Patterns

FC1 output & FC2 input channels converge to zero

Expert Layer Correlation

Patterns emerge in routed expert layers

Root Cause Analysis

NVFP4 vs BF16 Comparison

NVFP4 produces ~3× more zero-valued gradients

Underflow Mechanism

Small values ($<1e-12$) quantized to zero

Reversible Effect

Switching to BF16 returns zeros to baseline

Underflow Sources (Nemotron 3 Nano Study)

At 500B Tokens

High underflow in dgrad of FC2 → zeros in wgrad of FC1

At 750B Tokens

High underflow in fprop of FC1 → zeros in wgrad of FC2

2D weight quantization blocks span high & low magnitude channels causing differential underflow patterns



MXFP8 healing at 19T tokens tested but showed no downstream accuracy gains - final model uses NVFP4 through out



Pretraining Data Curation

Synthetic Code Concepts

91 concepts extracted

14M problems generated

15M final pairs

GPT-OSS 20B & 120B

Algorithmic Problems

Minimal prompts

LeetCode-style

0.2B tokens

Qwen3-235B + gpt-oss

Economics & Logic

Micro/macro/econometrics

Predicate/propositional logic

Truth tables

Qwen3-235B-Thinking

Synthetic Multiple Choice (MMLU-Style)

3.5M

Samples

1.6B

Tokens

+0.29

MMLU gain

+3.4

AIME gain

Pipeline: Bootstrap from MMLU auxiliary → Generate similar questions → DeepSeek-V3 solves with reasoning → Majority voting → Filter by agreement

Decontamination

- Exact match removal
- Embedding similarity >0.8 filtered
- Qwen3-Embedding-0.6

Cleaning Pipeline

- AST validation for code
- Import checking
- Format enforcement



Data Mixture & Training Phase

P1 Phase 1: Diversity (80%)

Tokens 20T

Focus Diversity & Coverage

Goal: Broad generalization across domains and tasks

P2 Phase 2: Quality (20%)

Tokens 5T

Focus High-Quality Sources

Goal: Refined performance on benchmarks

16 High-Level Data Categories

Web Crawl
5 quality groups

Math
Nemotron-CC-Math

Code
Nemotron-CC-Code

Wikipedia
High quality

Academic
Research papers

Crawl++
OWT, BigScience

Multilingual
Diverse languages

FinePDFs
Quality documents

SFT-Style Data

- general-sft
- stem-sft
- code-sft

Quality Estimation

Comparable quality → similar weights
Higher quality → greater weight

Nemotron-CC Taxonomy

crawl-medium → crawl-high
syn-crawl-medium-high → syn-crawl-high



Hyperparameters & Optimization

Learning Rate

Schedule: WSD
Warmup-Stable-Decay

Peak: 4.5×10^{-4}
@ 200B tokens

Min: 4.5×10^{-6}
Last 5T tokens

Optimizer

AdamW
Weight decay: 0.1

$\beta_1 = 0.9$
 $\beta_2 = 0.95$

Batch Configuration

Sequence Length
8,192

Batch Size
3,072 seqs

Tokens/Batch
~25.17M

MoE Configuration

Total Experts
512

Top-K
22

Router
Sigmoid + Bias

Load Balance
Aux-loss-free

MTP & Precision

MTP Loss Scale 0.3

Load Balance Update 10^{-3}

Aux Loss Coef 10^{-4}



Hybrid precision: BF16 for critical layers, NVFP4 for majority of network



Checkpoint Merging Strategy

The Challenge

Noisy Checkpoints

Stable LR phase exhibits step-to-step noise

Decay Run Cost

Dedicated decay runs require significant compute

Evaluation Overhead

Intermediate quality assessment is expensive

The Solution

Weighted Averaging

Merge over sliding window of checkpoints

Minus-Sqrt Decay

Emulation for merge coefficients

Compute Savings

~4T tokens (~16% of FLOP budget)

Results

2-4 pts

Improvement over trained checkpoint

125-500B

Merge window sizes evaluated

500B

Final model merge selected

Stable Phase

Merge consistently outperforms trained checkpoints by 2-4 points on 12 benchmarks

Decay Phase

Gap narrows substantially as trained checkpoints benefit from actual LR annealing



Long-Context Extension

LC-Phase Configuration

Learning Rate

4.5×10^{-6}

Constant

Batch Size

16

Global

Hardware

GB200 GPUs

Parallelism Strategy

Context Parallelism

64-way

Tensor Parallelism

2-way

Expert Parallelism

64-way

Two-Stage LC Training

Stage 1: 1M Context

Duration: 34B tokens

Context: 1,048,576 tokens

Data: 20% document QA, 80% Phase 2 blend

Stage 2: Alternating

Duration: 17B tokens

Context: Alternating 1M & 4K

Goal: Mitigate math benchmark impact

Document QA Data

Reused from Nemotron 2 & 3 Nano

20% of LC blend

Math Preservation



Alternating 4K sequences

Prevents degradation

Final Capability

Up to 1M context length



Strong long-context performance



CHAPTER 04

Post-Training Pipeline

Supervised fine-tuning and reinforcement learning for agentic capabilities





Supervised Fine-Tuning Approach

S1 Stage 1: Token-Level

Loss Function

$$L_{\text{tok}} = \sum \ell_t / |O_{\text{cl}}|$$

Global average over all output tokens

Configuration

256k packing, batch 64, LR 1e-5

Goal

Strong reasoning behavior

S2 Stage 2: Sample-Level

Loss Function

$$L_{\text{samp}} = (1/|B|) \sum (1/|O_{\text{cl}}|) \sum \ell_t$$

Per-conversation normalization

Configuration

512k packing, batch 32, LR 1e-5

Goal

Restore long-input-short-output

MTP During SFT

Shared-Weight Head

Same as pretraining

2 MTP Layers

Trained jointly

Loss Scale

0.3



Two-stage approach prevents degradation on long-input-short-output scenarios while retaining reasoning capabilities



SFT Data Categories

Software Engineering

SWE-Gym, R2E-Gym, SWE-rebench
OpenHands agent harness

Agentic Programming

20k queries → 15k tasks
+3k SWE + 10k web dev

Long Context

Multi-document reasoning
4-7 retrieval steps

Financial Reasoning

366k Q&A pairs
SecQue benchmark expansion

CUDA

100k samples
Kernel generation & optimization

Safety

Two-stage generation
Response policy framework

Search & Terminal

Search

Multi-hop web search

Terminal

84k samples

SQL & Tool Use

SQL

96.5k records

Tool Use

1.5M trajectories



Total: 7M+ samples | Significantly scaled agentic tasks vs Nemotron 3 Na

no



Reinforcement Learning Stage S

1

RLVR

21 environments
Math, code, STEM
Safety, agentic

2

SWE-RL

End-to-end SE
OpenHands agent
Multi-harness

3

RLHF

GenRM model
Principle-following
Helpsteer 3

4

MTP Healing

Frozen backbone
Train MTP heads
NLL loss

RLVR Environments (21 Total)

Math

Competitive + Python tool

Code

Competition-style

STEM

Scientific problems

Safety

Over-refusal + jailbreak

Agentic

Tool use + terminal

Reasoning Gym

Diverse reasoning

Low-Effort Reasoning

SFT Stage

2% of data from GPT-OSS low-effort

RL Stage

Starts at 2%, reduces to 1%

Reward Function

Correctness + token efficiency



Training on all environments simultaneously yields stable gains; single-environment training causes severe regressions



RL Infrastructure & Algorithms

Asynchronous GRPO

Decoupled Architecture

Training & inference on separate GPUs

Rollout Buffer

Continuous trajectory generation

In-Flight Updates

Weights pushed immediately

PivotRL for Agentic Tasks

Problem

SFT degrades OOD; RL is costly

Solution

Assistant-turn-level RL

Pivot Selection

Focus on uncertain turns

SWE-RL Infrastructure

Apptainer

Container isolation

OpenHands

Agent loop execution

Memory Watchdog

OOM prevention

Command Blocklist

Safety enforcement

NeMo Gym

Agent, model, resource servers

Ray + SLURM orchestration

NeMo RL



Training loop controller

Megatron-Core + vLLM

Scale

1k+ GPUs


21 environments, 37 datasets



CHAPTER 05

Quantization & Inference

Post-training quantization strategies for efficient deployment on
n Hopper and Blackwell





FP8 & NVFP4 Quantization

FP8 (W8A8) - Hopper

Calibration

256 samples, 65536 context

Quantized Layers

MoE GEMMs, Mamba Linear

KV Cache

FP8 for speedup

NVFP4 (W4A4) - Blackwell

Format

E2M1, 16-element blocks

Scaling

E4M3 per-block, FP32 global

Target

Prefill-heavy workloads

Hybrid NVFP4 Recipe

Weight Scales

Minimize MSE (offline)

Activation Scales

Max-based (runtime)

Layer Promotion

AutoQuantize (NAS)

AutoQuantize

Second-order Taylor approximation

Sensitivity + performance cost

Knapsack optimization

Results

99.8% median accuracy

<2 hours on 8x B200

512 samples, 4096 length



Mamba State Quantization

The Challenge

DRAM Bottleneck

SSM cache reads limit decoding speed

Error Accumulation

Recurrent: errors propagate over time

Verbosity Issue

FP16 → 40% verbosity increase

The Solution

FP16 + Stochastic Rounding

Zero-mean noise vs biased drift

Philox<5> PRNG

cuRAND support on Blackwell

PTX Instruction

Dedicated SR hardware

Error Accumulation Formula

$$h_{q,t} = h_t + \sum_{i=0}^t [\prod_{j=i+1}^t A_j] \cdot e_i$$

Quantization error from earlier steps propagates through subsequent recurrent transitions

Why Not INT16?

Requires block scales

Storage & loading overhead

Why Philox<5>?



Balances quality vs overhead

Maintains accuracy

Result

No verbosity increase



Accuracy maintained



CHAPTER 06

Performance Evaluation

Comprehensive benchmark results and comparisons with state
-of-the-art models





Base Model Evaluation Results



General Knowledge

MMLU	86.01
MMLU-Pro	75.65
GPQA	60.00

Math

GSM8K	90.67
MATH	84.84
AIME 2024	53.33

Code

HumanEval	79.40
MBPP	78.38
ARC-C	96.08

Long Context (RULER)

64K 93.2	128K 89.0	256K 86.2	512K 82.2	1M 74.4
-------------	--------------	--------------	--------------	------------

Comparison Summary

vs Ling-flash-Base-2.0	Superior on most
vs GLM-4.5-Air-Base	Superior on most
Long Context	Strong retention

Multilingual

MMLU Global Lite: 85.72
MGSM: 87.47

Reading

RACE: 91.00
HellaSwag: 88.97



Post-Trained Model Evaluation

Reasoning

AIME25: 90.21

HMMT: 94.73

GPQA: 79.23

Code

LiveCode: 81.19

SciCode: 42.05

HLE: 18.26

Agentic

SWE-B: 60.47

TauBench: 56.25

Terminal: 25.78

Chat & Context

IFBench: 72.56

Arena-Hard: 73.88

RULER 1M: 79.36

Key Highlights

SWE-Bench

60.47% (OpenHands) - competitive with baselines

Long Context

RULER 1M: 79.36% - strong retention at 1M tokens

Multilingual

MMLU-ProX: 86.67% - matches or outperforms

vs Qwen3.5-122B

Competitive on most benchmarks

Strong agentic performance

vs GPT-OSS-120B

Competitive or superior

Higher inference throughput



Quantized Model Accuracy Retention



Accuracy Comparison: BF16 vs FP8 vs NVFP4

MMLU-Pro	HMMT	GPQA	LiveCode	SciCode	HLE
83.73	94.73	79.23	78.69	42.05	18.26
83.63	94.38	79.36	78.44	41.38	17.42
83.33	95.36	79.42	78.44	40.83	17.42

Order: BF16 | FP8 | NVFP4

Terminal Bench

25.78 → 24.48

NVFP4 retention

TauBench V2

56.25 → 54.75

NVFP4 retention

RULER 1M

79.36 → 79.37

NVFP4 retention

Accuracy Retention

99.8%

Median accuracy vs BF16

Quantization Time

<2h

On 8× B200 GPUs



Inference Throughput Comparison

Test Configuration

Hardware
B200 GPUs

Framework
vLLM & TRT-LLM (best of two)

Sequence Length
8k input / 64k output

Throughput Gains

vs GPT-OSS-120B

2.2×

MXFP4 weights, MXFP8 activations

vs Qwen3.5-122B

7.5×

BF16 precision

Accuracy-Throughput Trade-off

IFBench
High
High

HMMT
High
High

SWE-B
High
High

HLE
High
High

Terminal
High
High

TauBench
High
High

RULER
High
High

Nemotron 3 Super achieves comparable accuracy with highest throughput

NVFP4 Advantage



Native Blackwell acceleration
Better than MXFP4

Hybrid Architecture

Mamba + Attention + MoE
Optimized for throughput

MTP Acceleration



Native speculative decoding
3.45 acceptance length



CHAPTER 07

Conclusion

Summary of contributions and open source impact





Key Contributions Summary

1

First NVFP4 Pretraining

Demonstrated stable training at 25T tokens in 4-bit precision, paving the way for efficient large-scale model training.

2

LatentMoE Architecture

Hardware-software co-design achieving better accuracy per FLOP and per parameter through latent space projections.

3

Shared-Weight MTP

Enables robust autoregressive speculative decoding with highest acceptance rates (3.45) on SPEED-Bench.

4

Hybrid Mamba-Attention MoE

Strategic global anchors support up to 1M context with efficient linear-time Mamba blocks.

5

Comprehensive Post-Training

21 RL environments with PivotRL for efficient agentic task training without OOD degradation.

6

Advanced Quantization

AutoQuantize with FP16+stochastic rounding for Mamba cache, achieving 99.8% accuracy retention.

7

Complete Open Source Release

All checkpoints (NVFP4, FP8, BF16), datasets (pretraining & post-training), and training recipes publicly available on HuggingFace and GitHub.



2.2× higher throughput than GPT-OSS-120B | 7.5× higher than Qwen3.5-12B

OPEN SOURCE RESOURCES

Build the Future of Efficient, Agentic AI



Model Checkpoints

- NVFP4 (W4A4)
- FP8 (W8A8)
- BF16 post-trained
- BF16 base
- GenRM for RLHF



Datasets

- Pretraining-Specialized-v1.1
- Super-Post-Training-Data
- Synthetic code concepts
- Agentic trajectories
- RL environments



Training Recipe

- NeMo framework
- NeMo Gym
- NeMo RL
- Model-Optimizer
- Complete pipeline



github.com/NVIDIA-NeMo/Nemotron

All resources available on HuggingFace

120B Total

12B Active

1M Context